

```
#####
```

```
# 2018/12
```

```
# Elasticsearch/Kibana: Version: 6.5.2
```

```
# reference: https://www.elastic.co/
```

```
# reference: Udemy: Complete Guide to Elasticsearch
```

```
#####
```

```
# node: a server that stores data and is part of a cluster; default name is UUID  
(Universally Unique Identifier)
```

```
# cluster: a collection of nodes (servers); default name is Elasticsearch
```

```
# document: each data item stored within a cluster is called document, being a basic  
unit of information that can be indexed; documents are stored within indices;  
documents are JSON objects
```

```
# index: an index is a collection of documents
```

```
# sharding divides indices into smaller pieces named shards; sharding enables you to  
distribute data across multiple nodes within a cluster; sharding also increases  
performance in cases where shards are distributed on multiple nodes because search  
queries can then be parallelized
```

```
# Elasticsearch natively support replication of your shards, meaning that shards are  
copied
```

```
# HTTP request methods
```

```
# GET: The GET method requests a representation of the specified resource. Requests  
using GET should only retrieve data.
```

```
# POST: The POST method is used to submit an entity to the specified resource, often  
causing a change in state or side effects on the server.
```

```
# PUT: The PUT method replaces all current representations of the target resource  
with the request payload.
```

```
# HTTP verb: GET, POST, PUT, DELETE
```

```
# <REST verb>/<index>/<type>/<API>      # <Type> will be removed in later/newer  
version of Elasticsearch
```

```
# create an index
```

```
PUT /examples_index_name/
```

```
# add/replace a document to an index
```

```
POST /examples_index_name/_doc
```

```
{  
  "my_field_1": "my_value",  
  "my_field_2": {  
    "my_field_3": true  
  },  
  "my_field_4": "2018-01-23T13:24:35+0800",  
  "my_field_5": null,  
  "my_field_6": 12.34  
}
```

```
# _id would be generated automatically, for example - "_id" : "_JLDuGoBjiKx_hYq32wY"
```

```
# add/replace a document with _id to an index
```

```

PUT /examples_index_name/_doc/<_id>
{
  "my_field_1": "my_value_1",
  "my_field_2": {
    "my_field_3": false
  },
  "my_field_6": 56.78
}

# retrieve a document by _id
GET /examples_index_name/_doc/<_id>

# using update API to update a document
POST /examples_index_name/_doc/<_id>/_update
{
  "doc": {"my_field_1": "my_value_update", "my_field_new": ["my_value_new"]}
}

# scripted update; use "script" property instead of "doc" property
POST /examples_index_name/_doc/<_id>/_update
{
  "script": "ctx._source.my_field_6 += 10"
}

# documents in Elasticsearch are immutable and cannot be changed

# upsert; if the document exist, then my_field_1 will be added by 5; if not exist,
my_field_1 will be 100
POST /examples_index_name/_doc/<_id>/_update
{
  "script": "ctx._source.my_field_1 += 5",
  "upsert": {
    "my_field_1": 100
  }
}

# delete a document
DELETE /examples_index_name/_doc/<_id>

# use _delete_by_query API to delete multiple documents
POST /examples_index_name/_delete_by_query
{
  "query": {
    "match": {
      "my_field_1": "my_value_1"
    }
  }
}

# bulk API can be used for add, update, and delete documents

```

```

# batch processing to add documents
POST /examples_index_name/_doc/_bulk
{"index": {"_id": "100"}}
{"my_field_1": "my_value_1"}
{"index": {"_id": "101"}}
{"my_field_1": "my_value_2"}

# batch processing to update/delete documents
POST /examples_index_name/_doc/_bulk
{"update": {"_id": "100"}}
{"doc": {"my_field_1": 1000}}
{"delete": {"_id": "101"}}

# import data with cURL
$ curl -H "Content-Type: application/json" -XPOST
"http://localhost:9200/examples_index_name/_doc/_bulk?pretty" --data-binary
"@my_data_file_1.jsonl" && curl -H "Content-Type: application/json" -XPOST
"http://localhost:9200/examples_index_name/_doc/_bulk?pretty" --data-binary
"@my_data_file_2.jsonl"

$ java -version

# use cat API for information about cluster, node, index, etc..
GET /_cat/health?v      # v for verbose
GET /_cat/nodes?v
GET /_cat/indices?v
GET /_cat/allocation?v  # how many shards are being allocated
GET /_cat/shards?v

# in Elasticsearch, mappings are used to define how documents and their fields
should be stored and indexed

GET /examples_index_name/_doc/_mapping

# each document has meta-data associated with it and they're called meta fields
# meta fields: _index, _id, _source, _field_names, _routing, _version, _meta
# field data types: core, complex, geo, specialized data types
# core: text, keyword (for filter/sort/aggregate), numeric, date, boolean
(true/false), binary, range
# complex: object, array, nested    # Lucene has no concept of inner objects
# geo: geo-point, geo-shape
# specialize: ip, completion, attachment

# create an index
PUT /examples
{
  "settings": {
    "index.number_of_shards": 1,
    "index.number_of_replicas": 0
  }
}

```

```

}

# create an index with mappings
PUT /examples_index_name/
{
  "mappings": {
    "my_default": {
      "dynamic": false,
      "properties": {
        "my_field_1": {
          "type": "integer"
        },
        "my_field_2": {
          "type": "boolean"
        },
        "latitude": {
          "type": "geo_point"
        }
      }
    }
  }
}

```

# mapping parameters: coerce (automatically cleaning up values), copy\_to, dynamic, properties, norms (used for relevance scores), format (for date fields), null\_value, fields

```

# add a mapping for new fields
PUT /examples_index_name/_doc/_mapping
{
  "properties": {
    "my_field_new_1": {
      "type": "date",
      "format": "yyyy/MM/dd HH:mm:ss||yyyy/MM/dd"
    },
    "my_field_new_2": {
      "type": "text",
      "fields": {
        "keyword": {
          "type": "keyword"
        }
      }
    }
  }
}

```

# existing mappings for fields cannot be updated

# analysis process involves tokenizing and normalizing a block of text  
 # text is tokenized into terms and terms are converted into lower case letters

(default behavior)

# results of analysis process are stored in inverted index

# inverted index is a mapping of a field's terms and which documents contain each term

# analyzer contains 1. character filter, 2. tokenizer, and 3. token filter

# character filter: manipulate text before tokenization; HTML strip character filter, mapping character filter, pattern replace

# tokenizer: split text into terms; word oriented tokenizers (standard, letter, lowercase, whitespace, UAX URL Email), partial word tokenizers (N-Gram, Edge N-Gram), structured text tokenizers (keyword, pattern, path)

# token filter: manipulate terms before adding them to an inverted index; standard, lowercase, uppercase, N-Gram, Edge N-Gram, stop, word delimiter, stemmer (going -> go), keyword marker, snowball, synonym, ASCII folding (résumé -> resume))

# built-in analyzer: standard, simple, stop, language, keyword, pattern, whitespace

# analyze API

POST \_analyze

```
{
  "tokenizer": "standard",
  "filter": ["lowercase", "unique", "asciifolding"],
  "text": "My@email.com mY.email my_email my-email résumé"
}
```

GET \_analyze

```
{
  "tokenizer": "letter",
  "filter": ["lowercase"],
  "text": "Brown.brown $brown fox $ brown @ dog https://www.google.com my@gmail.com"
}
```

GET \_analyze

```
{
  "tokenizer": "uax_url_email",
  "filter": ["lowercase"],
  "text": "Brown.brown $brown fox $ brown @ dog https://www.google.com my@gmail.com"
}
```

# define a custom analyzer which uses standard analyzer (configure built-in analyzers and token filters)

PUT /examples\_custom\_analyzer

```
{
  "settings": {
    "analysis": {
      "analyzer": {
        "my_english_stop": {
          "type": "standard",
          "stopwords": "_english_"
        }
      }
    }
  },
}
```

```

        "filter": {
            "my_stemmer": {
                "type": "stemmer",
                "name": "english"
            }
        }
    }
}

```

```

POST /examples_custom_analyzer/_analyze
{
  "analyzer": "my_english_stop",
  "text": "I'm in the mood of drinking semi-dry red wine!"
}

```

```

POST /examples_custom_analyzer/_analyze
{
  "tokenizer": "standard",
  "filter": ["my_stemmer"],
  "text": "I'm in the mood of drinking semi-dry red wine!"
}

```

```

PUT /examples_custom_analyzer_2
{
  "settings": {
    "analysis": {
      "analyzer": {
        "my_english_stop": {
          "type": "standard",
          "stopwords": "_english_"
        },
        "my_custom_analyzer": {
          "type": "custom",
          "tokenizer": "standard",
          "char_filter": [
            "html_strip"
          ],
          "filter": [
            "standard",
            "lowercase",
            "trim",
            "my_stemmer"
          ]
        }
      }
    },
    "filter": {
      "my_stemmer": {
        "type": "stemmer",
        "name": "english"
      }
    }
  }
}

```

```

    }
  }
}

```

POST /examples\_custom\_analyzer\_2/\_analyze

```

{
  "analyzer": "my_custom_analyzer",
  "text": "I'm in the mood for drinking <strong> semi-dry </strong> red wine! go
went fly flew"
}

```

# use analyzer in mapping

PUT /examples\_custom\_analyzer\_2/\_doc/\_mapping

```

{
  "properties": {
    "my_field_1": {
      "type": "text",
      "analyzer": "my_custom_analyzer"
    },
    "my_field_2": {
      "type": "text",
      "analyzer": "standard"
    },
    "my_fields_3": {
      "type": "text"
    }
  }
}

```

POST /examples\_custom\_analyzer\_2/\_doc/1

```

{
  "my_field_1": "drinking",
  "my_field_2": "drinking"
}

```

# query below will not hit any documents, since my\_field\_1 is stemmed to drink

GET /examples\_custom\_analyzer\_2/\_search

```

{
  "query": {
    "term": {
      "my_field_1": "drinking"
    }
  }
}

```

# add analyzers to existing indices (close index and then add analyzers, and then open index)

```
POST /examples_custom_analyzer_2/_close
```

```
PUT /examples_custom_analyzer_2/_settings
```

```
{
  "analysis": {
    "analyzer": {
      "my_french_stop": {
        "type": "standard",
        "stopwords": "_french_"
      }
    }
  }
}
```

```
POST /examples_custom_analyzer_2/_open
```

# the relevance scoring algorithm, OkAPI BM25, is currently used by Elasticsearch  
(it considers term frequency, inverse document frequency, field-length norm)

```
GET /examples_index_name/_search
```

```
{
  "explain": true,
  "query": {
    "match": {
      "my_field_1": "my_value"
    }
  }
}
```

```
GET /examples_index_name/_doc/<_id>/_explain
```

```
{
  "query": {
    "match": {
      "my_field_1": "my_value"
    }
  }
}
```

# query context affects relevance; filter context doesn't

# full-text queries

# term queries (enum, numbers, dates, etc..)

##### term queries #####

```
GET /examples_index_name/_doc/_search
```

```
{
  "query": {
    "term": {
```



```

        "my_field_1.keyword": "my_value"
    }
}

```

GET /examples\_index\_name/\_doc/\_search

```

{
  "query": {
    "term": {
      "my_field_7": {
        "value": false
      }
    }
  }
}

```

# search for multiple terms

GET /examples\_index\_name/\_doc/\_search

```

{
  "query": {
    "terms": {
      "my_field_7": [
        true,
        false
      ]
    }
  }
}

```

# note: term/terms & my\_field\_7 is boolean type

# retrieve documents based on IDs

GET /examples\_index\_name/\_doc/\_search

```

{
  "query": {
    "terms": {
      "_id": ["1", "2", "3"]
    }
  }
}

```

# use range query for number or date

# example: date math & round time

# 2 year and one day before and round date to month

# round current time to month and subtract a year

GET /examples\_index\_name/\_doc/\_search

```

{
  "query": {
    "range": {
      "my_field_4": {

```

```

        "gte": "01-01-2013||-2y-1d/M",
        "lte": "now/M-1y",
        "format": "dd-MM-yyyy"
    }
}

# match document with non-null value
GET /examples_index_name/_doc/_search
{
  "query": {
    "exists": {
      "field": "my_field"
    }
  }
}

# match based on prefix
GET /examples_index_name/_doc/_search
{
  "query": {
    "prefix": {
      "my_field.keyword": "veget"
    }
  }
}

# search with wildcards (performance is slow)
# *: match any characters sequence including no characters
# "vege*able": start with vege and end with able
# ?: match any single character
# "vege?able"
GET /examples_index_name/_doc/_search
{
  "query": {
    "wildcard": {
      "my_field.keyword": "vege*able"
    }
  }
}

# search with regular expressions
# Elasticsearch uses Lucene regular expression engine
GET /examples_index_name/_doc/_search
{
  "query": {
    "regexp": {
      "my_field.keyword": "veget[a-zA-Z]+ble"
    }
  }
}

```

```

    }

}
##### term queries #####

```

```

##### full text queries #####
# match query operates on terms

```

```

GET /examples_index_name/_doc/_search
{
  "query": {
    "match": {
      "my_field_1": {
        "query": "my_value",
        "operator": "and"
      }
    }
  }
}

```

```

# terms in specific order (exact phrase)
GET /examples_index_name/_doc/_search
{
  "query": {
    "match_phrase": {
      "my_field_1": "my_value"
    }
  }
}

```

```

# search multiple fields
GET /examples_index_name/_doc/_search
{
  "query": {
    "multi_match": {
      "query": "my_value",
      "fields": ["my_field_1", "my_field_2"]
    }
  }
}

```

```

##### full text queries #####

```

```

##### compound queries #####
GET /examples_index_name/_doc/_search
{

```

```

"query": {
  "bool": {
    "must": [
      {
        "match": {
          "my_field_1": "my_value"
        }
      },
      {
        "range": {
          "my_field_6": {
            "lte": 15
          }
        }
      }
    ]
  }
}

```

# move range query to filter object to improve efficiency  
 # if there is a match for "should", score will be boosted  
 # if there is a "must", "should" is optional and having a match for "should" will boost score.  
 # if only has "should" for example no "must", then records have to match "should" and "should" becomes required.

GET /examples\_index\_name/\_doc/\_search

```

{
  "query": {
    "bool": {
      "must": [
        {
          "match": {
            "my_field_1": "my_value"
          }
        }
      ],
      "must_not": [
        {
          "match": {
            "my_field_1": "my_value_2"
          }
        }
      ],
      "should": [
        {
          "match": {
            "my_field_1": "my_value_3"
          }
        }
      ]
    }
  }
}

```



```

    ],
    "filter": [
      {
        "range": {
          "my_field_6": {
            "lte": 15,
            "_name": "my_identifier_4"
          }
        }
      },
      {
        "term": {
          "my_field_1.keyword": "my_value"
        }
      }
    ]
  }
}

```

# the two queries example 1 and 2 are the same  
 # one of them should match

# example 1:

```

GET /examples_index_name/_doc/_search
{
  "query": {
    "match": {
      "my_field_1": "my_value_1 my_value"
    }
  }
}

```

# example 2:

```

GET /examples_index_name/_doc/_search
{
  "query": {
    "bool": {
      "should": [
        {
          "term": {
            "my_field_1": "my_value_1"
          }
        },
        {
          "term": {
            "my_field_1": "my_value"
          }
        }
      ]
    }
  }
}

```

```

    }
  }
}

```

# in elasticsearch, you generally denormalize data for best performance instead of mapping document relationships  
 # don't use Elasticsearch as a primary data store  
 # Elasticsearch only supports simple joins (joins are expensive)  
 ##### compound queries #####

##### nested #####  
 # query type: nested

```

PUT /examples_nested
{
  "mappings": {
    "_doc": {
      "properties": {
        "name": {
          "type": "text"
        },
        "employees": {
          "type": "nested"
        }
      }
    }
  }
}

```

# if employees aren't defined as nested datatype, association among Eric, 39, M would loss

```

POST /examples_nested/_doc/1
{
  "name": "Development",
  "employees": [
    {
      "name": "Eric",
      "age": 39,
      "gender": "M"
    },
    {
      "name": "John",
      "age": 20,
      "gender": "M"
    }
  ]
}

```

POST /examples\_nested/\_doc/2

```
{
  "name": "HHR",
  "employees": [
    {
      "name": "Maria",
      "age": 30,
      "gender": "F"
    },
    {
      "name": "Ryan",
      "age": 25,
      "gender": "M"
    }
  ]
}
```

GET /examples\_nested/\_search

```
{
  "query": {
    "nested": {
      "path": "employees",
      "query": {
        "term": {
          "employees.gender.keyword": "M"
        }
      }
    }
  }
}
```

# note that result is 2 instead of 3; might need to loop through each object to get total counts of 3

# number of department has M is 2

# number of employees with gender = M is 3

# nested inner hit

GET /examples\_nested/\_search?filter\_path=hits.hits.inner\_hits.employees.hits.total

```
{
  "_source": false,
  "query": {
    "nested": {
      "path": "employees",
      "inner_hits": {},
      "query": {
        "term": {
          "employees.gender.keyword": "M"
        }
      }
    }
  }
}
```



```

}
##### nested #####

##### join #####
# define parent/child relationship for my_department/my_employee
PUT /examples_join
{
  "mappings": {
    "_doc": {
      "properties": {
        "my_join_field": {
          "type": "join",
          "relations": {
            "my_parent_department": "my_child_employee"
          }
        }
      }
    }
  }
}

PUT /examples_join/_doc/2
{
  "my_name": "R&D",
  "my_join_field": "my_parent_department"
}

PUT /examples_join/_doc/3
{
  "my_name": "HHR",
  "my_join_field": "my_parent_department"
}

# parent/child documents must be stored on the same shard
# routing: in which shard a document with a given ID is stored
# using parent routing id = 2 in an example
PUT /examples_join/_doc/4?routing=2
{
  "my_name": "Bo R&D",
  "my_age": 10,
  "my_gender": "F",
  "my_join_field": {
    "name": "my_child_employee",
    "parent": 2
  }
}

PUT /examples_join/_doc/5?routing=2

```

```
{
  "my_name": "John R&D",
  "my_age": 40,
  "my_gender": "F",
  "my_join_field": {
    "name": "my_child_employee",
    "parent": 2
  }
}
```

PUT /examples\_join/\_doc/6?routing=3

```
{
  "my_name": "Bill HHR",
  "my_age": 45,
  "my_gender": "F",
  "my_join_field": {
    "name": "my_child_employee",
    "parent": 3
  }
}
```

PUT /examples\_join/\_doc/7?routing=3

```
{
  "my_name": "Joe HHR",
  "my_age": 48,
  "my_gender": "M",
  "my_join_field": {
    "name": "my_child_employee",
    "parent": 3
  }
}
```

# query based on parent id

GET /examples\_join/\_search

```
{
  "query": {
    "parent_id": {
      "type": "my_child_employee",
      "id": 2
    }
  }
}
```

# query child document by parent

GET /examples\_join/\_search

```
{
  "query": {
    "has_parent": {
      "parent_type": "my_parent_department",
      "score": true,

```

```

        "query": {
            "term": {
                "my_name.keyword": "R&D"
            }
        }
    }
}

```

# query parent document by child

GET /examples\_join/\_search

```

{
    "query": {
        "has_child": {
            "type": "my_child_employee",
            "score_mode": "sum",
            "min_children": 2,
            "max_children": 10,
            "query": {
                "bool": {
                    "must": [
                        {
                            "range": {
                                "my_age": {
                                    "gte": 10,
                                    "lte": 46
                                }
                            }
                        }
                    ],
                    "should": [
                        {
                            "term": {
                                "my_gender.keyword": "M"
                            }
                        }
                    ]
                }
            }
        }
    }
}

```

# define multiple-level relations

PUT /examples\_join\_multi\_level

```

{
    "mappings": {
        "_doc": {
            "properties": {
                "my_join_field": {

```



```
}  
  
PUT /examples_join_multi_level/_doc/6?routing=4  
{
```

```
  "my_name": "Bill HHR",  
  "my_join_field": {  
    "name": "my_child_employee",  
    "parent": 5  
  }  
}
```

```
GET /examples_join_multi_level/_search  
{
```

```
  "query": {  
    "has_child": {  
      "type": "my_parent_department",  
      "query": {  
        "has_child": {  
          "type": "my_child_employee",  
          "query": {  
            "term": {  
              "my_name.keyword": "Bill HHR"  
            }  
          }  
        }  
      }  
    }  
  }  
}
```

```
# parent/child inner hits  
GET /examples_join/_search
```

```
{  
  "query": {  
    "has_parent": {  
      "parent_type": "my_parent_department",  
      "score": true,  
      "inner_hits": {},  
      "query": {  
        "term": {  
          "my_name.keyword": "R&D"  
        }  
      }  
    }  
  }  
}
```

```
GET /examples_join/_search
```

```
{
  "query": {
    "has_child": {
      "type": "my_child_employee",
      "score_mode": "sum",
      "min_children": 2,
      "max_children": 10,
      "inner_hits": {},
      "query": {
        "bool": {
          "must": [
            {
              "range": {
                "my_age": {
                  "gte": 10,
                  "lte": 46
                }
              }
            }
          ],
          "should": [
            {
              "term": {
                "my_gender.keyword": "M"
              }
            }
          ]
        }
      }
    }
  }
}
```

# terms lookup mechanism (fetch a term from documents)

# join limitations:

- # documents must be stored within the same index
- # parent/child documents must be on the same shard
- # only one join field per index; a join field can have as many relations as you want; new relations can be added after creating the index
- # child relations can only be added to existing parents
- # document can only have one parent (employee only belongs to one department)

##### join #####

# specify result formats (yaml/json)

GET /examples\_index\_name/\_doc/\_search?format=yaml

```
{
  "query": {
```

```
        "match_all": {}
    }
}
```

# filter source

```
GET /examples_index_name/_doc/_search
{
  "_source": false,
  "query": {
    "match_all": {}
  }
}
```

# only my\_fields under \_source is returned

```
GET /examples_index_name/_doc/_search
{
  "_source": ["my_field_1", "my_field_2"],
  "query": {
    "match_all": {}
  }
}
```

```
GET /examples_index_name/_doc/_search
{
  "_source": {
    "includes": ["my_field_1", "my_field_2"],
    "excludes": "my_field_2"
  },
  "query": {
    "match_all": {}
  }
}
```

# limit result size; hits.total still has correct counts of results

# default size is 10

```
GET /examples_index_name/_doc/_search?size=2
{
  "query": {
    "match_all": {}
  }
}
```

```
GET /examples_index_name/_doc/_search
{
  "size": 2,
  "query": {
    "match_all": {}
  }
}
```

# use offset to specify the number of matches to skip before returning the matches  
GET /examples\_index\_name/\_doc/\_search

```
{
  "from": 2,
  "query": {
    "match_all": {}
  }
}
```

# can use "from" and "size" for pagination

# sort

GET /examples\_index\_name/\_search

```
{
  "_source": false,
  "query": {
    "match_all": {}
  },
  "sort": [
    "my_field_6"
  ]
}
```

# sort key includes a number, date will be # of millisecond since the epoch (unix timestamp \* 1000)

GET /examples\_index\_name/\_search

```
{
  "query": {
    "match_all": {}
  },
  "sort": [
    {"my_field_4": "desc"},
    {"my_field_6": "asc"}
  ]
}
```

# sort by multi-value fields

GET /examples\_index\_name/\_search

```
{
  "query": {
    "match_all": {}
  },
  "sort": [
    {
      "my_field_6": {
        "order": "desc",
        "mode": "avg"
      }
    }
  ]
}
```



```
}
```

```
##### filter #####
```

```
# filter is more efficient
```

```
# typically use filter for number & date ranges and term query for keyword
```

```
GET /examples_index_name/_search
```

```
{
  "query": {
    "bool": {
      "must": [
        {
          "match_all": {}
        }
      ],
      "filter": [
        {
          "range": {
            "my_field_6": {
              "lte": 2.5
            }
          }
        }
      ]
    }
  }
}
```

```
##### filter #####
```

```
##### aggregations #####
```

```
# metric aggregations
```

```
# do aggregation on all documents
```

```
GET /examples_index_name/_search
```

```
{
  "size": 0,
  "aggs": {
    "my_avg": {
      "avg": {
        "field": "my_field_6"
      }
    },
    "my_sum": {
      "sum": {
        "field": "my_field_6"
      }
    }
  }
}
```

```

}

# count unique values
# use my_field.keyword because Fielddata is disabled on text fields by default or
# enable fielddata on a text field ("fielddata": true)
# index a field as both "text" and "keyword"
# cardinality aggregation produces an approximate number
# the precision_threshold options allows to trade memory for accuracy
GET /examples_index_name/_search

```

```

{
  "size": 0,
  "aggs": {
    "my_unique_field": {
      "cardinality": {
        "field": "my_field_1.keyword"
      }
    }
  }
}

```

```

GET /examples_index_name/_search
{
  "size": 0,
  "aggs": {
    "my_unique_field": {
      "terms": {
        "field": "my_field_1.keyword"
      }
    }
  }
}

```

# value\_count aggregation: count number of values that the aggregation is based on  
 (# of values that are extracted from the aggregation documents)

```

GET /examples_index_name/_search
{
  "size": 0,
  "aggs": {
    "my_value_count": {
      "value_count": {
        "field": "my_field_6"
      }
    }
  }
}

```

# multi-value aggregation: stats

```

GET /examples_index_name/_search
{
  "size": 0,

```

```

    "aggs": {
      "my_stats": {
        "stats": {
          "field": "my_field_6"
        }
      }
    }
  }
}

```

# bucket aggregations

# terms aggregation: build buckets for each unique value

# elasticsearch only returns top unique terms so some terms won't appear in results;  
 "sum\_other\_doc\_count" key is the sum of document counts which aren't part of the results

# "missing" is for documents contain null values or don't have my\_field at all; user can name the bucket "my N/A" for "missing"

# "min\_doc\_count" key specifies minimum number of documents a bucket needs to contain to be present in results (default is 1)

# "doc\_count" is approximate and not always accurate

# "\_term" is a special key which allows users to refer to buckets' keys

# "doc\_count\_error\_upper\_bound" means the max possible document counts for a term isn't part of the final results

GET /examples\_index\_name/\_search

```

{
  "size": 0,
  "aggs": {
    "my_terms": {
      "terms": {
        "field": "my_field_1.keyword",
        "missing": "my N/A",
        "min_doc_count": 0,
        "order": {
          "_term": "asc"
        }
      }
    }
  }
}

```

# nested aggregation (sub-aggregation)

GET /examples\_index\_name/\_search

```

{
  "size": 0,
  "query": {
    "range": {
      "my_field_6": {
        "lte": 100
      }
    }
  },

```

```

    "aggs": {
      "my_terms": {
        "terms": {
          "field": "my_field_1.keyword"
        },
        "aggs": {
          "my_terms_sub": {
            "stats": {
              "field": "my_field_6"
            }
          }
        }
      }
    }
  }
}

```

# filter out document in aggregation

GET /examples\_index\_name/\_search

```

{
  "size": 0,
  "aggs": {
    "my_range": {
      "filter": {
        "range": {
          "my_field_6": {
            "lte": 5000
          }
        }
      },
      "aggs": {
        "my_avg": {
          "avg": {
            "field": "my_field_6"
          }
        }
      }
    }
  }
}

```

# define bucket rules with filters

GET /examples\_index\_name/\_search

```

{
  "size": 0,
  "aggs": {
    "my_filter": {
      "filters": {
        "filters": {
          "my_bucket_rule_1": {
            "match": {

```



```

"size": 0,
"aggs": {
  "my_range": {
    "date_range": {
      "field": "my_field_4",
      "format": "yyyy-MM-dd",
      "keyed": true,
      "ranges": [
        {
          "from": "2017-07-07",
          "to": "2017-07-07||+6M",
          "key": "my_bucket_name_1"
        },
        {
          "from": "2017-07-07||+6M",
          "to": "2017-07-07||+1y",
          "key": "my_bucket_name_2"
        }
      ]
    },
    "aggs": {
      "my_stats": {
        "stats": {
          "field": "my_field_6"
        }
      }
    }
  }
}
}

```

# histogram - bucket results using values (min/max/specified interval)  
 GET /examples\_index\_name/\_search

```

{
  "size": 0,
  "aggs": {
    "my_histogram": {
      "histogram": {
        "field": "my_field_6",
        "interval": 10,
        "min_doc_count": 0,
        "extended_bounds": {
          "min": 0,
          "max": 100
        }
      }
    }
  }
}

```

# date\_histogram

GET /examples\_index\_name/\_search

```
{
  "size": 0,
  "aggs": {
    "my_date_histogram": {
      "date_histogram": {
        "field": "my_field_4",
        "interval": "month"
      }
    }
  }
}
```

# global aggregation isn't influenced by the search query

GET /examples\_index\_name/\_search

```
{
  "size": 0,
  "query": {
    "range": {
      "my_field_6": {
        "gte": 20
      }
    }
  },
  "aggs": {
    "my_aggs_range": {
      "stats": {
        "field": "my_field_6"
      }
    },
    "my_aggs_all": {
      "global": {},
      "aggs": {
        "my_global_stats": {
          "stats": {
            "field": "my_field_6"
          }
        }
      }
    }
  }
}
```

# missing field values

POST /examples\_missing/\_doc/1

```
{
  "my_field_1": 100,
  "my_field_2": "text"
}
```

```
POST /examples_missing/_doc/2
```

```
{
  "my_field_1": 200,
  "my_field_2": null
}
```

```
GET /examples_missing/_search
```

```
{
  "size": 0,
  "aggs": {
    "my_missing_aggs": {
      "missing": {
        "field": "my_field_2.keyword"
      },
      "aggs": {
        "my_missing_sum": {
          "sum": {
            "field": "my_field_1"
          }
        }
      }
    }
  }
}
```

```
# aggregating nested objects
```

```
# it is important to include the path of the object while working on nested query
```

```
GET /examples_nested/_search
```

```
{
  "size": 0,
  "aggs": {
    "my_nested_min": {
      "nested": {
        "path": "employees"
      },
      "aggs": {
        "my_min": {
          "min": {
            "field": "employees.age"
          }
        }
      }
    }
  }
}
```

```
##### aggregations #####
```



##### improve search results #####

# add more data

POST /examples\_index\_name/\_doc/\_bulk

```
{"index": {"_id": "100"}}
{"my_field_1": "dog brown", "my_field_8": "vegetable"}
{"index": {"_id": "101"}}
{"my_field_1": "brown dog", "my_field_8": "veget"}
{"index": {"_id": "102"}}
{"my_field_1": "brown a dog", "my_field_8": "vegetables"}
{"index": {"_id": "103"}}
{"my_field_1": "doggy", "my_field_8": "vegetable"}
```

# proximity searches

# slop 2: phrase can be reversed

GET /examples\_index\_name/\_search

```
{
  "query": {
    "match_phrase": {
      "my_field": {
        "query": "my_value",
        "slop": 1
      }
    }
  }
}
```

# match query will match documents with my\_value\_0 or my\_value\_1 ("vegetable" or "veget")

# "match\_phrase" under "should" is optional and it boosts scores

GET /examples\_index\_name/\_search

```
{
  "query": {
    "bool": {
      "must": [
        {
          "match": {
            "my_field_8": {
              "query": "vegetable veget"
            }
          }
        }
      ],
      "should": [
        {
          "match_phrase": {
            "my_field_1": {
              "query": "brown dog",
              "slop": 2
            }
          }
        }
      ]
    }
  }
}
```

```

    }
  ]
}
}
}

```

```

# fuzzy match query for typo; levenshtein distance (edit distance); maximum edit
distance: 2
# automatic fuzziness:
# term length 1-2 uses maximum edit distance 0
# term length 3-5 uses maximum edit distance 1
# term length >5 uses maximum edit distance 2
# lobster & oyster are two levenshtein distance apart; majority of human typos is
edit distance 1

```

```
GET /examples_index_name/_search
```

```

{
  "query": {
    "match": {
      "my_field_8": {
        "query": "vegetable",
        "fuzziness": "auto"
      }
    }
  }
}

```

```

# transposition (live v.s. lvie) is 1 edit distance apart
# or to set fuzzy_transpositions to be false and edit distance is 2

```

```
GET /examples_index_name/_search
```

```

{
  "query": {
    "match": {
      "my_field_8": {
        "query": "live",
        "fuzziness": 1,
        "fuzzy_transpositions": true
      }
    }
  }
}

```

```

# fuzzy query (term-level query) isn't analyzed
# match query with fuzziness option (full-text query): analyzed; preferred method
than fuzzy query

```

```
GET /examples_index_name/_search
```

```

{
  "query": {
    "fuzzy": {
      "my_field_8": {
        "value": "VEGETABLE",

```

```

        "fuzziness": "auto"
    }
}

```

GET /examples\_index\_name/\_search

```

{
  "query": {
    "match": {
      "my_field_8": {
        "query": "VEGETABLE",
        "fuzziness": "auto"
      }
    }
  }
}

```

# synonyms

# "lowercase" filter is loaded before "my\_synonym\_test" filter

PUT /examples\_synonyms\_index

```

{
  "settings": {
    "analysis": {
      "filter": {
        "my_synonym_test": {
          "type": "synonym",
          "synonyms_path": "analysis/my_synonym_test.txt"
        }
      },
      "analyzer": {
        "my_analyzer": {
          "tokenizer": "standard",
          "filter": [
            "lowercase",
            "my_synonym_test"
          ]
        }
      }
    }
  },
  "mappings": {
    "_doc": {
      "properties": {
        "my_description": {
          "type": "text",
          "analyzer": "my_analyzer"
        }
      }
    }
  }
}

```

```
}  
}
```

##### my\_synonym\_test.txt (config/analysis/my\_synonym\_test.txt) ##### all nodes have to have this file

```
# This is a comment  
awful => terrible    # awful will be replaced by terrible (note: term query isn't  
analyzed so won't work)  
awesome => great, super  
elasticsearch, logstash, kibana => elk    # all lower cases because "lowercase"  
filter is loaded before "my_synonym_test" filter  
weird, strange    # both two terms are placed in the same position
```

##### my\_synonym\_test.txt #####

```
# use analyzer to test synonyms  
POST /examples_synonyms_index/_analyze  
{  
  "analyzer": "my_analyzer",  
  "text": "awesome"  
}
```

```
POST /examples_synonyms_index/_analyze  
{  
  "analyzer": "my_analyzer",  
  "text": "ELASTICSEARCH"  
}
```

```
POST /examples_synonyms_index/_analyze  
{  
  "analyzer": "my_analyzer",  
  "text": "ELASTICSEARCH is weird"  
}
```

```
POST /examples_synonyms_index/_doc  
{  
  "my_description": "Elasticsearch is awesome, but can also seem weird sometimes."  
}
```

```
# there is a match even though there is no "great" in my_description;  
"my_synonym_test" works
```

```
GET /examples_synonyms_index/_doc/_search  
{  
  "query": {  
    "match": {  
      "my_description": "great"  
    }  
  }  
}
```

```
GET /examples_synonyms_index/_doc/_search
{
  "query": {
    "match": {
      "my_description": "awesome"
    }
  }
}
```

# note: if there are documents indexed before adding synonyms, update by query API needs to be executed so all documents will be re-index

```
POST /examples_synonyms_index/_update_by_query
```

# highlight matches in fields

```
POST /examples_highlight_index/_doc/1
```

```
{
  "my_description": "Elasticsearch is a search engine based on the Lucene library."
}
```

```
GET /examples_highlight_index/_doc/_search
```

```
{
  "query": {
    "match": {
      "my_description": "Lucene library"
    }
  },
  "highlight": {
    "pre_tags": ["<my_tag>"],
    "post_tags": ["</my_tag>"],
    "fields": {
      "my_description": {}
    }
  }
}
```

# stemming

```
PUT /examples_stemming_index
```

```
{
  "settings": {
    "analysis": {
      "filter": {
        "my_synonym_test": {
          "type": "synonym",
          "synonyms": [
            "firm => company",
            "love, enjoy"
          ]
        }
      }
    }
  }
}
```

```

        "my_stemmer_test" : {
            "type" : "stemmer",
            "name" : "english"
        }
    },
    "analyzer": {
        "my_analyzer": {
            "tokenizer": "standard",
            "filter": [
                "lowercase",
                "my_synonym_test",
                "my_stemmer_test"
            ]
        }
    }
},
"mappings": {
    "_doc": {
        "properties": {
            "my_description": {
                "type": "text",
                "analyzer": "my_analyzer"
            }
        }
    }
}
}

```

POST /examples\_stemming\_index/\_doc/1

```

{
  "my_description": "I love working for my firm!"
}

```

# search enjoy, love, work, or working

GET /examples\_stemming\_index/\_doc/\_search

```

{
  "query": {
    "match": {
      "my_description": "enjoy work love working"
    }
  },
  "highlight": {
    "fields": {
      "my_description": {}
    }
  }
}

```

##### improve search results #####

##### other notes #####

# document contains fields

# text search in Elasticsearch

# score is based on relevance (\_score)

# find "field" contains "text or phrase"

# "must": all conditions have to be true for results to return

# "should" - can weight criteria differently using boost

# match, match with fuzziness, match\_phrase, match\_phrase\_prefix, match\_phrase with slop

# boost score

GET /examples\_index\_name/\_search

```
{
  "query": {
    "bool": {
      "should": [
        {
          "match_phrase": {
            "my_field_1": "brown dog"
          }
        },
        {
          "match_phrase": {
            "my_field_1": {
              "query": "a brown dog",
              "boost": 3
            }
          }
        }
      ]
    }
  }
}
```

"""

GET \_search

```
{
  "min_score": 0,
  "size": 250,
  "query": {
    "match_all": {}
  },
  "highlight": {
    "fields": {
      "my_field": {}
    }
  },
}
```

```

"sort": [
  {"my_field" : "desc"}
],
"aggs": {
  "my_statistics": {
    "stats": {
      "field" : "my_field.keyword",
      "size": 10
    }
  },
  "my_unique_count": {
    "terms": {
      "field": "my_field"
    }
  },
  "my_doc_count": {
    "range": {
      "field": "my_field",
      "ranges": [
        {
          "key" : "0-10000",
          "from" : 0,
          "to" : 10
        }
      ]
    }
  }
}
}

```

# SQL in Elasticsearch

POST /\_xpack/sql?format=txt

```

{
  "query": "SELECT * FROM examples_index_name WHERE my_field_1 = 'brown dog'"
}

```

# list all stored scripts

GET \_cluster/state/metadata?pretty&filter\_path=\*\*.stored\_scripts

# delete a stored script

DELETE \_scripts/my\_stored\_script

# register a stored script

POST \_scripts/my\_stored\_script

```

{
  "script": {
    "lang": "mustache",
    "source": {
      "query": {

```



```

        "match": {
            "my_field_1": "{{var_my_field_1}}"
        }
    }
}

```

```

# retrieve a stored script
GET _scripts/my_stored_script

```

```

# use a stored script
GET /examples_index_name/_search/template
{
    "id": "my_stored_script",
    "params": {
        "var_my_field_1": "my_value"
    }
}

```

```

GET /_msearch
{ "index": "examples_index_name" }
{ "query": { "match": { "my_field_1": "my_value" } } }
{ "index": "examples_index_name" }
{ "query": { "match": { "my_field_1": "dog" } } }

```

```

$ bin/logstash -e "input { stdin { } } output { stdout { } }"

```

```

$ cd \logstash\bin

```

```

$ logstash -f my_logstash.conf

```

```

##### my_logstash.conf #####

```

```

input {
    file {
        path => "my_input_data.csv"
        id => "my_plugin_id"
        start_position => "beginning"
        sincedb_path => "/dev/null"
    }
}
filter {
    csv {
        separator => ","
        columns => [ "my_column_1", "my_column_2", "my_column_3" ]
    }
    mutate {convert => ["my_column_2","integer"]}
    mutate {convert => ["my_column_3","float"]}
}
output {
    elasticsearch {
        hosts => "localhost"
        index => "examples_logstash_index"
    }
}

```

```

    }
    stdout {codec => dots}
}
##### my_logstash.conf #####

# examples of dates which Elasticsearch can recognize by default
"2019-05-01" or "2019/05/01 13:15:30"

# geo-distance
PUT /examples_lat_long
{
  "mappings": {
    "_doc": {
      "properties": {
        "my_address": {
          "properties": {
            "latlong": {
              "type": "geo_point"
            }
          }
        }
      }
    }
  }
}

PUT /examples_lat_long/_doc/1
{
  "my_address" : {
    "latlong" : {
      "lat" : 32.936461,
      "lon" : -117.234071
    }
  }
}

PUT /examples_lat_long/_doc/2
{
  "my_address" : {
    "latlong" : {
      "lat" : 32.936529,
      "lon" : -117.233178
    }
  }
}

GET /examples_lat_long/_search
{
  "query": {
    "bool" : {
      "must" : {

```



```
    }  
  }  
)
```

```
my_es_client.count(  
    index="examples_index_name",  
    body={  
        "query": {  
            "match_all": {}  
        }  
    }  
)
```

```
my_es_client.search_template(  
    index="examples_index_name",  
    body={  
        "id": "my_stored_script",  
        "params": {  
            "var_my_field_1": "my_value"  
        }  
    }  
)
```

```
##### Python #####
```